

The Software Toolworks®

15233 VENTURA BOULEVARD SUITE 1118, SHERMAN OAKS, CALIFORNIA 91403 (213) 986-4885

TOOLWORKS C/80
Version 3.0
July 1983
Walt Bilofsky

Table of Contents

1. INTRODUCTION.....	3
2. LEARNING C; OTHER REFERENCES.....	5
3. FOR THE EXPERIENCED C USER.....	6
4. THE C/80 DISTRIBUTION DISKS.....	7
5. AN EXAMPLE.....	9
6. RUNNING THE COMPILER.....	10
6.1. Compiler Switches.....	10
6.2. Changing Compiler Defaults.....	12
7. C/80 LANGUAGE SUMMARY.....	13
7.1. Variables.....	13
7.2. Data Types.....	15
7.3. Pointers.....	16
7.4. Structures.....	16
7.5. Storage Classes.....	17
7.6. Constants.....	17
7.7. Operators and Expressions.....	18
7.8. Statements.....	19
7.9. Conclusion.....	19
8. IMPLEMENTATION AND MACHINE DEPENDENCIES.....	20
9. RUNTIME AND I/O LIBRARY.....	23
9.1. Files and Devices.....	23
9.2. Commands.....	24
9.3. I/O Redirection.....	24
9.4. Interrupting a Program.....	25
9.5. Basic I/O Library Routines.....	26
9.6. Formatted Input and Output.....	28
9.7. More Storage and I/O Routines.....	30
9.8. Arithmetic and Number String Functions.....	30
9.9. String Manipulation.....	31
9.10. CP/M System Calls.....	32
9.11. Random Access File I/O.....	32
9.12. Program Chaining; Wildcards.....	33

10. USING C/80 WITH MACRO-80 OR RMAC.....	34
10.1. Assembling and Loading.....	34
10.2. ROMable Code.....	35
10.3. Making Libraries.....	35
11. MULTIPLE COMPILES USING AS.....	36
12. RUNTIME TRACE AND EXECUTION PROFILE.....	37
13. THE AS ASSEMBLER.....	38
14. UPPER CASE SOURCE FILES.....	39
15. TRICKS AND INTERNALS.....	40
15.1. Global Arrays.....	40
15.2. Assembly Language Linkage.....	40
15.3. Declarations for Efficient Code.....	40
15.4. Defining Globals in Header Files.....	40
15.5. I/O Buffers.....	41
15.6. Crash on Exit.....	41
16. COMPILER ERROR MESSAGES.....	42
INDEX.....	46

Acknowledgements.

Ron Cain's contribution in providing a simple, public domain compiler for a minimal C subset is well known and widely appreciated. Jim Gillogly wrote and maintained `printf`. The initial version of `seek` and `HDOS exec` were contributed by Al Bolduc. `CP/M exec` was written by Robert Wesson. Floats and longs were implemented by Herve Tireford of SGS-ATES, Geneva. Grant Gustafson contributed `scanf`. Tyler Sperry donated the bibliography. Version 3.0 reflects contributions and suggestions from, and fixes of bugs reported by, Jim and Gary Gilbreath, Bruce Wampler, and many other users.

Copyright (c) 1981, 1982, 1983 Walter Bilofsky. Sale of this software conveys a license for its use on a single computer at a time, owned or operated by the purchaser. Copying this software or documentation by any means whatsoever for any other purpose is expressly prohibited. C/80 and Toolworks are trademarks of The Software Toolworks. CP/M is a registered trademark of Digital Research.

No License Fees: COM or ABS files incorporating executable versions of the C/80 runtime library may be copied or distributed without restriction or fee, although credit to C/80 is appreciated. With that exception, none of the files on the distribution disk, including C, ASM, REL and COM files, may be copied or distributed in any other form, except as provided in the preceding paragraph, under penalty of law.

TOOLWORKS C/80
Version 3.0
July 1983
Walt Bilofsky

1. INTRODUCTION

C/80 is a compiler for the C programming language, running under the CP/M and HDOS operating systems. It requires a minimum of 56K of memory. The compiler produces an assembly language text file which is turned into an executable object program by the AS absolute assembler, which is included. Optionally, C/80 can produce output for Microsoft's Macro-80 or Digital Research's RMAC relocatable assembler.

The reference manual for C/80 is The C Programming Language by Brian Kernighan and Dennis Ritchie. Section 2 tells where you can obtain this book, and lists other useful books on C.

Purchasers of inexpensive C compilers have come to expect that they will lack many important language features, or will have non-standard variations which make C programs less portable. C/80 Version 3.0 is one of the better compilers in this respect. It supports all of the language features described in The C Programming Language, with the following exceptions:

- o float and long data (available option; see below)
- o double data type
- o typedef
- o Arguments to #define macros
- o Bit fields
- o #line
- o Declarations within nested blocks

C/80 Version 3.0 does support structures, statics, initialization, casts, compile time evaluation of constant expressions -- in short, all other C language features. A few language features have restrictions on their use; see Section 8 for a complete list of the exceptions and implementation dependencies.

Float and long data types are not supported in the basic C/80 compiler, but can be added with the optional C/80 MATHPAK. This allows us to offer the basic compiler at an extremely affordable price, while still making floats and longs available for those who wish them.

- o UNIX-style I/O redirection and command line expansion.
- o Conventional C I/O and string library
- o Formatted and random access file I/O
- o Dynamic storage allocation
- o Runtime execution profile facility
- o Selectable Macro-80 or RMAC compatibility
- o In-line assembly language
- o Includes absolute assembler

Note: This document describes C/80 implementations for both CP/M and HDOS. Where program names, devices, etc., differ for the two systems, the CP/M names will be used, with the HDOS equivalent in brackets, [like this].

1. The first step is to identify the problem or goal. This involves understanding the current situation and what needs to be achieved.

080 MATHEMATICS

1. 凡在本行開辦之各項業務，
 2. 均應遵守本行所定之各項規章，
 3. 並應隨時注意本行所定之各項規章，
 4. 如有違反者，本行將依法究辦。

2. LEARNING C; OTHER REFERENCES

A summary and brief description of the C/80 language appears below (Section 7). However, for a detailed introduction to C, the beginner will need more than this manual provides. There are several books on C, suitable for readers at various stages of expertise. Many are available at local bookstores, or can be readily ordered there. Prices are subject to change.

The C Programming Language, Brian Kernighan & Dennis Ritchie, 1978 Prentice-Hall, Englewood Cliffs, NJ., 228 pages. The definitive work (i.e., it defines the language), but not for beginners. You will probably want it eventually if not immediately.

C Programming Guide, Jack Purdum, 1983 Que Corporation, Indianapolis, IN., \$17.95, 250 pages. Aimed at CP/M and beginners. Assumes some minimal experience with BASIC or another programming language.

Learning to Program in C, Thomas Plum, 1983 Plum Hall, Cardiff, NJ., \$29.95, 350 pages. Good for beginners; great for those who have a PDP-11 or an interest in what the computer does with the code.

The C Primer, Les Hancock & Morris Krieger, 1982 McGraw-Hill (Byte Books), New York, NY., \$14.95, 235 pages. Good for those starting out, but examples often assume UNIX environment. Not as complete as Plum or as useful to C/80 users as Purdum, but good.

The C Puzzle Book, Alan Feuer, 1982 Prentice-Hall, Englewood Cliffs, NJ., \$14.95, 173 pages. A book of exercises, grouped by topic (operators, pointers, etc.). Intended to be used with another book.

C Notes, C. T. Zahn, 1979 Yourdon Press, New York, NY., \$18.50, 100 pages. More a manual than a tutorial; more organized than Kernighan & Ritchie. Not a primer.

An Introduction to C, Bruce Hunter, available late 1983, Sybex. For beginning or intermediate C programmers; deals with existing C compilers in the context of CP/M and UNIX.

All these books are paperbacks. If not available through your local bookstore, they may be ordered from:

Opamp Technical Books, Inc.
1033 N. Sycamore Ave.
Los Angeles, CA 90038
(213) 464-4322

Note: In reading these books, it is important to keep in mind at least the major features of C which are not supported by C/80. Section 8 below lists the differences between C/80 and the language described in The C Programming Language.

3. FOR THE EXPERIENCED C USER

Users who know C will find C/80 quite familiar and easy to use. It supports data types `char` and `int`, and with the optional `MATHPAK`, `long` and `float`. C/80 has full C pointers, arrays and structures, all C control statements, all operators, and most preprocessor functions. The preprocessor allows in-line assembly language code.

C/80 programs use the conventional command line (`main(argc,argv)`). The runtime library provides many of the capabilities of the standard C I/O library file handling routines, and implements UNIX-style I/O redirection in the command line. The `command` library function is provided to perform UNIX-style command line wildcard expansion.

There are a few differences between C/80 and full C. The one difference which may cause hard to detect problems in converting C programs written for other compilers is that functions may not be called with a different number of arguments than specified in the function definition. The I/O is closer to Version 7 `stdio` than it appears; see the end of Section 8 for the correspondence.

C/80 program source files are normally prepared using the full upper and lower case ASCII character set. Users with upper case only terminals should refer to Section 14.

Users of earlier versions of C/80 will notice the following changes, improvements and new features in C/80 3.0:

- o Expanded runtime library (Section 9)
 - o ROMable code (Section 10.2)
 - o RMAC compatibility (Section 10)
- o Menu configurable compiler (Section 6.2)
 - o \ at end of line for continuation
- o -Q switch for `#define` (Section 6.1)
- o `#ifneeded` for selective compile (Section 8)
- o True `alloc/free` (Section 9.7)
 - o Command line wildcard expansion (Section 9.12)
- o CP/M files are written in 128 byte records
- o File 0 is always the terminal (Section 9.5)

4. THE C/80 DISTRIBUTION DISKS

C/80 comes on one or two disks, containing the following files:

Files You Will Need to Get Started:

C.COM or C.ABS The C/80 compiler.

AS.COM [or .ABS] An absolute 8080 assembler which assembles C/80 output files. See Section 13.

CLIBRARY.ASM The basic C/80 library. When the AS assembler is used, this file is automatically included in C/80 programs. Usually it should reside on A: [SY0: on HDOS] at assembly time.

PRINTF.C The C/80 formatted output routines (see Section 9.6). May be incorporated into a program using `#include printf.c`, which should appear before any use of `printf`.

Example Programs:

HELLO.C A sample C/80 program; the first program in The C Programming Language. See Section 5.

TAB.C A sample C/80 program which copies a file, replacing blanks by tabs wherever this might result in a savings of space.

CMP.C A sample C/80 program which compares two files.

TREE.C A sample C/80 program showing the use of structures.

Files Containing Library Functions:

TPRINTF.C A smaller version of `printf` without left justification or precision.

PRINTF.H A header file for defining `printf` in source files where the entire `printf.c` or `tprintf.c` source file is not included.

SCANF.C The C/80 formatted input routines (see Section 9.6).

SCANF.H A header file for defining `scanf` in source files where the entire `scanf.c` source file is not included.

STDLIB.C Library of C and assembly language routines for string and character manipulation, system calls, and storage allocation. See Sections 9.7 - 9.10.

SEEK.C Routines for random access file I/O (not for CP/M 1.4 and earlier). See Section 9.11.

EXEC.C Routine to chain to another .COM [.ABS] program. See Section 9.12.

COMMAND.C Routine to expand file wildcards in command line. See Section 9.12.

Files Used with Macro-80 or RMAC:

CLIBRARY.REL A relocatable version of CLIBRARY, for use with Microsoft's Link-80 or LINK from Digital Research.

STDLIB.REL A relocatable library version of STDLIB.

Files for Advanced Users:

CCONFIG.COM [or .ABS] Program for changing defaults in C.COM. See Section 6.2.

CLIBIO.C Commented source code for the I/O and system dependent portions of CLIBRARY.

CPROF.C The runtime execution profile library. Compile to create CPROF.REL or CPROF.ASM. The latter file is automatically included in assemblies of C/80 programs compiled with the -p switch (see Section 12). It must reside on A: [SY0: on HDOS] when such programs are assembled.

CTRACE.ASM An alternate runtime execution profile library which traces each routine call and return. See Section 12.

5. AN EXAMPLE

This section describes how to assemble and run a C/80 program. The example shown is for the CP/M operating system. [Under HDOS, the procedure is identical, except that HDOS uses the ">" prompt instead of "A>", and the Heath assembler ASM is used, with the command "asm hello=hello".]

First a source file, called HELLO.C, must be prepared. This can be done using any text editor. (Or you can use the file HELLO.C on the C/80 distribution disk.) The program on the source file should look like this (more or less):

```
#include "printf.c"
main() {
    printf("Hello, world!\n");
}
```

Files PRINTF.C and CLIBRARY.ASM should be copied to A: [SY0: on HDOS]. Then HELLO can be compiled and run by the following steps. Characters which the computer types are underlined in this example; the other characters are typed by the user.

A> c hello

C/80 Compiler 3.0 (7/27/83)

A> as hello

8080 AS 2.3 (7/15/83)

A> hello

Hello, world!

A>

44A.80A9T1

6. RUNNING THE COMPILER

The simplest way to compile a C/80 program is to give the command

```
c filename
```

This takes the source file FILENAME.C and produces an assembly file FILENAME.ASM. Of course, any file name can be used instead of FILENAME.

In order to create an assembly file with a different name or on a different device, the command looks like

```
c d:outfile.mac=b:infile
```

```
[c dk1:outfile.mac=syl:infile on HDOS]
```

If no extension is specified, the defaults are .ASM for the output file (except .MAC for Macro-80 files), and .C for the input.

After compiling the source program, use AS to assemble the assembly language file into an executable program. Note that all C programs will include the C runtime library file, CLIBRARY.ASM, which must reside on A: [SY0: on HDOS]. The syntax of the AS command is the same as for C, unless a listing file is desired (Section 13).

6.1. Compiler Switches.

The C command may include "switches" to select compiler options. The switches consist of a -, a letter (upper and lower case are synonymous), and sometimes a numeric value (represented below by N). Switches can be separate (-t -s400) or strung together (-ts400t).

Example: To compile file FOO.C, producing file B:FOO.ASM, including the source text as comments in the assembly file, and allocating space for 400 symbols, use the command

```
c -t -s400 b:foo=foo
```

Once you have determined the switch values you usually use, you can change the compiler to default to those values; see Section 6.2.

6.2.

The switches are:

Assembler Output Format

-t Include the source program text as comments in the assembly language file.

- m (or -m1). Generate Macro-80 assembler output.
- m2 Generate RMAC assembler output. See Section 10.
- m0 Generate AS assembler output (the default).
- a Do not generate the relocation assembly directives CSEG and DSEG in the assembly language file. This option is only valid when the -m switch is selected. This will help if the linker runs out of space. See Section 10.1.
- p Generate a runtime profile for the program being compiled.
- vB: Find CLIBRARY.ASM (and other ASM library files) on the specified disk instead of A:. Only used with AS assembler. See Section 12.

Compilation

- qV=S Specify a **#define** in command line. This switch has the same effect as the preprocessor statement **"#define V S"**, where V is an identifier and S is a string. Note that the command line is always translated to upper case, so V and S will be taken as upper case. -qV defines V as the null string. If the -d switch is also used, it must appear before the -q switch in the command line.
- eN Set page length for error message pauses. Normally, compiler will pause after 24 lines of error messages. -e0 eliminates pauses entirely.

Table Storage Allocation

- sN Allocate N entries in the symbol table. 3/4 of the symbol table entries are used for globals, and the remainder for local variables. The symbol table uses about 16 bytes per entry.
- cN Allocates N bytes for the string constant table. This table stores all string constants in the program (but see the -k switch below).
- dN Allocates N bytes for the **#define** table. This table stores all **#define** macros and strings.
- wN Allocates N slots for the switch/case table. The size of this table determines the maximum number of cases in a switch statement or in nested switch statements.
- rN Allocates N bytes for the structure table. This table stores the information from structure declarations.

Multiple Assembly Options (with AS)

- lN Begin generating internal labels at number N (default 1). This provides a method of compiling several C source files into separate assembly files which may then be assembled together. See Section 11.
- g Do not reserve storage for globals. This is equivalent to preceding each global declaration with the `extern` keyword. This may be useful in multiple file compilations; see Section 11.

Data Initialization Options

- k Normally, the compiler outputs string constants at the end of the compilation. The `-k` switch dumps string constants after each function. This can greatly reduce the amount of memory used, allowing larger programs to be compiled, but duplicate strings may not be detected (see next paragraph).
- f Normally, the compiler tries not to duplicate strings which can be overlapped. In particular, two identical string constants will point to the same location. This may cause problems in a program which alters the contents of a string constant or an initialized character pointer. (Initialized character arrays are not affected.) The `-f` switch turns off the string overlap feature, and insures that each string is stored separately.
- z C generally initializes all static and global storage to zeros. Since this can create a very large intermediate assembly language file, C/80 only zeros uninitialized arrays shorter than 256 bytes. The `-z` switch causes all statics and globals to be initialized to zero, regardless of size.

6.2. Changing Compiler Defaults.

You may change most of the default values of the compiler switches in the previous section by running the CCONFIG program supplied on your C/80 distribution disk. This program will patch new values for the defaults into the file C.COM. In particular, you can specify defaults for the size of most tables, the drive for CLIBRARY.ASM, screen size for error pause, choice of assembler, and generation of CSEG and DSEG.

The program is menu driven and mostly self explanatory. You should remember to include a drive identifier when specifying a C.COM file on other than the current drive.

If you configure a switch to be on, using it in the command line will turn it off. For example, if you configure C.COM to default to Macro-80, then the `-m` switch in the command line will specify the AS assembler.

7. C/80 LANGUAGE SUMMARY

This section contains a language summary in tabular form, followed by a concise explanation of the major C language features. Its aim is to convey the basics of C/80 to a programmer with some feel for computer languages. It is not intended to replace The C Programming Language as an exhaustive reference, and where there is a conflict, the latter is the authority (except for implementation differences listed in Section 8).

Long and float data types are available only with the optional C/80 MATHPAK; see Section 1.2.

(1) Data Types:

Types: char, int, unsigned
long, float (optional)

Declarations:

```
char c, *pc, ac[], ac2[n], **x[m][n]
int i, *pi, ai[], aci[n];
long j, *pj, bi[], bci[n];
float k, *pk, ci[], cci[n];
extern char/int ...
static char/int ...
auto char/int ...
register char/int ...
initialization:
```

```
char/int v = constant;
```

```
char/int a[] = {c,c,...};
```

Structure Declarations: See "Structures" below.

(2) Primaries:

constant:

```
decimal number
octal number beginning with 0
hex number beginning with 0x or 0X
character constant 'c'
string "abc"
```

variable

```
address[expression]
```

```
function(arg1,...,argn) (n >= 0)
```

```
structure.element
```

```
ptr_to_structure->element
```

(3) Expressions:

Unary operators:

```
-      minus
*      contents of
&      address of
++     increment (pre- or postfix)
--     decrement (pre- or postfix)
!      truth valued not
~      bitwise not
(type) type is any type (e.g., char*);
      forces type of following expression
sizeof nr. bytes in type or expression
```

Binary operators:

*	/	arithmetic operators
%		modulo
+	-	arithmetic operators
>>	<<	right, left shift
<	<=	less than, less than or equal
>	>=	greater than, greater than or equal
==	!=	equal, not equal (0 or 1 valued)
&	^	bitwise and, or, xor
&&		truth valued and, or
?:		if-then-else expression
=		assignment operator
+=	-=	arithmetic assignment operators
*=	/=	
&=	=	
^=	>>=	
<<=		
,		comma (expression sequencing)

(4) Statements

```

expression;
if (expression) statement;
if (expression) statement; else statement;
for (expression; expression; expression) statement
while (expression) statement;
do statement while (expression);
switch (expression) {
    case: statement; ...
    default: statement; }
break;
continue;
return;
return expression;
goto label;
label: statement;
{ statement; ...; statement; }
;      (null statement)

```

(5) Function Definitions (functions are fully recursive)

```

fname(arg1,...,argn)
    int/char argi,*argj, ... ;
    { statements; }

```

(6) Preprocessor Functions

/* comments */	
#define name string	Replace name by string throughout text
#undef name	Erase definition
#include "filename"	Inserts filename at that point.
or #include <filename>	
#ifdef name	Generate following code if name is #defined or not #defined, or if expression nonzero, or if an undefined global, respectively.
#ifndef name	
#if expression	
#ifneed name,...,name	
#else	Reverse conditional generation
#endif	End conditional generation

#asm	Begin assembly language
#endasm	End assembly language
#UPPER	Convert upper to lower case

7.1. Variables.

Variable names consist of letters, numbers, and the character "_". The first character must be a letter. Upper and lower case letters are allowed and are different (except globals when Macro-80 or RMAC are used); usually lower case letters are used. Variable names may be any length, but any letters after the first seven are ignored (six for globals with Macro-80 or RMAC).

Each variable has a type, a scope, and a storage class.

The scope of a variable determines the portion of the source program within which the variable is known. The three possible scopes are local, global, and external. Local variables are those declared at the beginning of a function body; they are known only inside that function and the same names can be used in other functions. Global variables are those declared outside a function body; they are known in all functions from the declaration to the end of the file. In addition, the extern declaration may be used to reference variables in the C/80 library or on other files in a multiple file compilation.

7.2. Data Types

C/80 contains two basic data types: int, which is a 16 bit signed integer (range -32768 to 32767); and char, which is an 8 bit signed integer (range -128 to 127). Chars are often used to store characters of text. Unsigned ints are also supported (range 0 to 65535).

The optional MATHPAK adds the data types long, which is a 32 bit signed integer (range -2,147,483,648 to 2,147,483,647); and single-precision float, which is a 32-bit quantity, with about 7 decimal digits of precision, in the range e-38 to e+38.

Integers, characters, longs, and floats are the basic components of the more complex data types: arrays, pointers, and structures. The following declaration declares an int, a doubly dimensioned int array, an array of pointers to ints, and a function returning pointers to ints.

```
int i, array[30][10], *ptr[5], *f();
```

An array is similar to a BASIC or FORTRAN array; it simply consists of consecutive pieces of memory, each large enough to contain a char, int or pointer. An n long array may be subscripted from 0 to n-1 only.

C does not contain a separate string data type; strings are stored as arrays of chars. By convention, a byte containing 0 is used to terminate a string.

7.3. Pointers

The concept of pointers is essential to the C language. A pointer is simply an address. Thus, in C/80, pointers are unsigned 16 bit numbers, similar to unsigned ints. They are used to step efficiently through arrays, where other languages might use subscripts and an index variable, and to pass (addresses of) large data structures as function arguments. An indication of how to use pointers is given in Section 7.7.

7.4. Structures

Structures are a useful way to organize data. An example of a structure declaration is:

```
struct tree {  
    char value[5];  
    struct tree *left,*right; }  
    forest[50],*ptree;
```

This declares three kinds of things: a structure type called **tree**, structure elements called **value**, **left** and **right**, and variables: an array of structures, called **forest**, and a pointer to objects of type **struct tree**, called **ptree**. Each object of type structure is like an array. But whereas an array contains a number of pieces of data all of the same type, a structure contains pieces of data called elements, which may be declared to be of different types. In the declaration above, objects of type **struct tree** are declared to contain a 5-long character array, and two pointers to things of type **struct tree**.

In a structure declaration you can omit the variables, the structure name, or (once the structure type has been declared) the { ... }.

A structure element can be referred to by the operators **->** or **::**

```
forest[i].value  
ptree->left
```

These elements can be treated just the same as variables. Use **.** to refer to fields of variables which are structures, and **->** to refer to fields of things which point to structures.

Structures and their use are a complicated topic and can only be touched on here. The C Programming Language and the other books mentioned in Section 2 contain discussions and examples which will be helpful.

7.5. Storage Classes

Storage classes determine how a variable is stored in memory. The storage classes in C/80 are:

```
static
auto
register
extern
```

Statics and **externs** may be either local to a function or global; **auto** and **registers** can only be local.

Statics are simply memory locations. **Auto** variables are stored on the pushdown stack. Local **statics** (declared within a function) are preserved between calls of that function, whereas **autos** are not. When a function is called recursively, a new **auto** is created local to that call of the function; **statics** are not. Local variables default to **auto** but may be declared **static**. (C/80 actually uses static storage for local 16 bit variables and saves the values so that this is transparent to the user but more efficient than using the stack. The explicit **auto** declaration may be used to override this and force the variables to be located on the stack.) **Register** variables are stored efficiently, but are otherwise the same as **autos**.

Variables declared outside a function can be either global or external (see Section 7.1). An explicit **static** declaration declares variables known to the end of the source file, but not in other source files. **Externs** are globals which are declared in another source module. Omitting the declarator **extern** or **static** defines a global which can be referred to as an **extern** from other modules.

7.6. Constants.

A decimal constant consists of a string of decimal digits. A constant beginning with '0' (e.g., 0177) is interpreted as an octal number. A constant beginning with '0x' or '0X' is a hex constant.

C/80 computes 16 bit constant expressions at compile time. Wherever an integer constant is required (such as the dimension in an array declaration), you can use an expression containing only integer constants.

C/80 also contains string, and character constants. Characters are any printing character, or

```

\t      for tab
\n      for newline (end of line)
\r      for carriage return
\b      for backspace
\f      for form feed
\\      for \
\'      for '
\"      for "
\123    for the octal value 123
        (or any other value)

```

A single character constant is written 'c'. A string constant is written "ccc...". A string constant is stored as a 0-terminated array of `chars`. Useful things to do with string constants include assigning them to `char` pointers, and passing them as arguments to functions.

7.7. Operators and Expressions.

The operators in C are shown in the table at the beginning of Section 7. They appear in the approximate order in which they are performed during expression evaluation; e.g., `/` is performed before `+` is performed before `&`.

Certain operators are peculiar to C. The unary operator `*` takes a pointer and yields the contents of the location it points to. The operator `&` takes an object, which must have an address, and yields a pointer to that address. Thus, for any expression `A` which has an address, the value of `*&A` is the same as `A`.

The operators `++` and `--` can appear either before or after a variable. They cause the variable to be incremented (`++`) or decremented (`--`) by 1. The value of the expression is the variable either before (`V++`) or after (`++V`) the operation. For example, if `p` is a pointer to `char`, then `*p++` increments `p`, but applies the `*` operation to the value of `p++`, which is the value of `p` before being incremented. This leads to the following sort of code, which is common in C:

```

p = "Any old string\n";
while (*p) putchar(*p++);

```

This outputs the string by calling `putchar()` with each character in it, until the 0 byte terminating the string is encountered.

When a pointer is incremented or decremented, its value changes not by 1 but by 1 object. Thus incrementing a pointer to an `int` moves it to point to the next following `int`; its actual value increases by 2. A pointer to a structure is incremented by the size of the structure, in bytes.

Truth values in C are either zero (false) or nonzero (true). Truth-valued operators (`==`, `>`, `&&`, etc.) return 1 for true. A useful expression in C is

```

expr ? tvalue : fvalue

```

whose value is tvalue if expr is nonzero, and fvalue if expr is zero.

7.8. Statements.

The table at the beginning of Section 7 lists the statements in C. Anywhere that C allows a single statement, it will also accept a compound statement of the form

```
{ statement; statement; ...; }
```

The iterative statements in C all have simple equivalent definitions:

```
for (e1; e2; e3) statement;
```

```
means      e1; L: if (e2) { statement; e3; goto L; }
```

```
while (e) statement;
```

```
means      L: if (e) { statement; goto L; }
```

```
do statement; while (e);
```

```
means      L: statement; if (e) goto L;
```

```
switch (e) { case c1: s1; ... case cn: sn; default: s; }
```

```
means      if (e == c1) s1;
            :
            if (e == cn) sn;
            s;
```

The **switch** statement is not exactly equivalent to what is shown, however. The expression *e* is actually evaluated only once. The **case** values *c1*, ..., *cn* must be constants. And the **default** case may be eliminated, in which case no **case** is executed if the value *e* is different from the values *c1*, ..., *cn*.

The statement **break** jumps out of the smallest **for**, **while**, or **switch** containing it. The statement **continue** begins the next iteration of the smallest **for** or **while** containing it.

7.9. Conclusion.

This short a summary can not begin to present all the details of the C language. In order to learn more, you can look at the sample C programs provided on the C/80 distribution disk and read the books referenced in Section 2.

8. IMPLEMENTATION AND MACHINE DEPENDENCIES

The reference manual for C/80 is The C Programming Language (see Section 2). In using that book, it must be kept in mind that some features of C and its runtime library are not present in C/80, or differ from the description in the book. Those omissions and differences are listed here. First the unimplemented features and major restrictions are listed, followed by a detailed listing of all differences. Section numbers refer to the C Reference Manual in Appendix A of The C Programming Language.

Unimplemented Features

Float, double, entry and typedef keywords (2.3)
Long and float constants (2.4) and arithmetic.
(Float and long may be added with the
optional C/80 MATHPAK.)
Typedef (8.1, 8.8)
Bit fields (8.5)
#line (12.4)

Major Restricted Features

Function calls (7.1) must have the same number of
arguments as the called function definition.
Blocks (9.2): declarations are allowed only at start
of a function.
#define (12.1): arguments are not allowed.

Implementation Dependencies and Difference List

2 Lexical conventions

Blanks in the middle of multiple character operators (e.g.,
=*) are not allowed.

2.2 Names

The first seven characters of a name are significant. If
Macro-80 or RMAC are used, global symbols are restricted to 6
characters, and upper and lower case are the same in globals.

2.3 Keywords

Float, double, entry and typedef keywords are not
recognized.

2.4 Constants

Long and floating constants are not implemented.

2.6 Hardware characteristics

Char is 8 bits. Int, short and long are 16 bits.

7.1 Primary expressions.

Function calls must have the same number of arguments as the called function definition.

8.1 Storage class specifiers

Register variables must be at most 16 bits long. They are stored in static memory for fast access, and are saved on function entry and exit. C/80 allows any number of register variables, and the & operator may be applied to them, but such code may not be portable.

Auto variables are stored on the stack. Local variables default to register if they are 16 bits long, but the `auto` declaration can override this. Other local variables, and all arguments, default to `auto`, but the register declaration can be used to override this (except for variables longer than 16 bits). All this is transparent to the user.

The scope of an `extern` declaration is the remainder of the source file, even if the declaration is within a function definition.

`Typedef` is not recognized.

8.5 Structure and union declarations

Bit fields are not implemented.

8.6 Initialization

Only static and global variables may be initialized. Only objects smaller than 256 bytes are defaulted to 0; the `-z` compiler switch removes this restriction. (See also Sect. 15 below.)

Type declarations can not be nested. About the only restriction this imposes is that `sizeof` a type name may not be used in a dimension in a declaration.

8.8 Typedef

`Typedef` is not implemented.

9.2 Compound statement, or block

Declarations are allowed only at the beginning of a function body.

12 Compiler control lines

In-line assembly language is supported by the `#asm` and `#endasm` directives.

A new conditional compilation directive is supplied:

`#ifneed namel,...,namen`

The code following, up to the matching `#endif`, will be compiled if one or more of the names `namel,...,namen` are externs not yet defined in the compilation. This directive can be used to create C `#include` files containing a number of library routines, each of which is compiled only if it has been called.

12.1 Token replacement

`#Define` is not applied recursively, and arguments to macros are not allowed.

12.4 Line control

`#line` is not implemented.

15 Constant expressions

`?:` and `&` are not allowed in constant expressions, except that `&` may be the first character in an initializer.

17 Anachronisms

All forms listed are recognized, except that initializers which lack an `=` and start with `(name...` will not compile.

Printf and scanf:

These functions (and the variants, `fprintf`, etc.) cope with a variable number of arguments through use of a `#define` kludge which redefines `printf`. This requires that either `#include "printf.c"` or `#include "printf.h"` be placed before the first use of `printf`, and similarly for `scanf.h` and `scanf`. Also, a use of `scanf` must be enclosed in parentheses in order to return the correct value. See Section 9.6.

I/O and Runtime Library:

Many of the basic library functions and treatments of files described in the manual are supported, although the format is not always identical. `Getchar`, `putchar`, `getc`, `putc`, `fopen`, `fclose`, `seek`, `ftell` and `exec` are provided. EOF has the value `-1`; `NULL` is `0`.

The older Version 6 convention of `fin` and `fout` is followed instead of `stdin`, `stdout` and `stderr`. However, compatibility may be maintained with most implementations by taking the following definitions:

```
#define FILE int
extern int fin,fout;
#define stdin fin
#define stdout fout
#define stderr 0
```

I/O redirection is provided. The I/O and runtime library is described more fully in the next section.

9. RUNTIME AND I/O LIBRARY

C/80 is supplied with a runtime library, which provides convenient access to files and other devices in a manner generally consistent with accepted C conventions.

The library is divided into several sections, contained on different files. Some are provided as C/80 source files, some as ASM and/or REL files.

CLIBRARY	Basic I/O routines; always loaded
STDLIB	Additional I/O, string manipulation, memory allocation and CP/M system calls
PRINTF,SCANF	Formatted I/O
SEEK	Random access file I/O
EXEC	"Chaining" to another program
COMMAND	Command line file name wildcard expansion
CPROF,CTRACE	Runtime execution trace and profile
CLIBIO	Source for I/O portion of CLIBRARY

9.1. Files and Devices:

NOTE: The following information applies only to use of C/80 on the CP/M operating system, and should be ignored by users of HDOS.

Under CP/M, the C/80 library recognizes the logical device driver names CON:, RDR:, PUN: and LST: as legal file names.

In doing I/O to CON:, C/80 normally uses line at a time mode. This is true whether CON: has been accessed explicitly by opening file "CON:", or as the default device for getchar and putchar. If you need to use character at a time console I/O, declare

extern char Cmode;

and set Cmode to zero. (It is initially set to 1; other values produce undefined results.)

When reading from CON:, in addition to the control characters which are interpreted by CP/M (see the description of functions 1 and 10 in the CP/M 2.2 Interface Guide), C/80 interprets ctrl-B as an interrupt. Moreover, when doing I/O to any of the four logical devices, C/80 maps CR into '\n' on input and '\n' into CR-LF on output. Calling fopen to open the device in binary mode has no effect on this (although it does for file I/O); the only way to avoid these mappings is not to go through the C/80 library.

When one of these devices is closed, the end of file character is written to it. The char array IOeof[4] holds the end of file character for CON:, LST:, RDR:, and PUN: respectively.

[Under HDOS, any legal file or device name may be used whenever a file name is called for. Examples of legal names are FOO, SY1:FILE.DAT, or LP:.]

Upper and lower case letters are legal and synonymous in file names. On CP/M, the user is responsible for insuring that characters which confuse the CCP, such as '.' and ':', are not used.

9.2. Commands:

C/80 programs begin execution by calling the routine `main`, which the user must provide. `Main` should start off with the declaration

```
main(argc,argv)
char *argv[];
```

When `main` is called, `argc` will be the number of elements in `argv`, and `argv` will be an array of pointers to the strings which appear on the command line (except for `argv[0]`, which may not contain anything useful.) `Argv[argc]` is always -1. For example, if a C/80 program named `progl` is run with the command

```
progl a b foodle
```

then `argc` is 4, and `argv` contains pointers to the strings "PROGL" (or nonsense), "A", "B" and "FOODLE" (e.g., `argv[3]` is "FOODLE"). An argument containing spaces and/or tabs may be enclosed in either single or double quotes.

Due to limitations in both the HDOS and CP/M systems, lower case letters in the argument line are passed to `main` as the upper case equivalent.

9.3. I/O Redirection:

Many C programs do input and output a character at a time, taking input from the standard input using the routine `getchar`, and writing to the standard output using the routines `putchar`, or `printf`. The standard input and standard output are both initially the terminal. However, files or other devices can also be used as the standard input and/or output.

This is implemented by two global extern ints, `fin` and `fout`. These variables define the standard input and output, respectively. They are usually set to 0 before `main` is called, and I/O is done to and from the terminal. However, a program may open a file for reading or writing, and set `fin` or `fout` to the file's channel number. This will cause subsequent I/O to be done to the file rather than to the terminal. A device may be used instead of a file.

Normally, an output file must be closed explicitly by a program before exiting. However, the I/O library always closes `fout` when a program is terminated (including abnormal termination by `ctrl-C` under HDOS and `ctrl-B` under CP/M).

Fin and **fout** may also be redefined in the command line when the program is run. For example, the command

```
progl a b <b:infile >lst:
```

will run **progl** with arguments "A" and "B". **Getchar()** will read characters from file "B:INFILE", and **putchar()** will write characters to the printer. The **>** and **<** arguments are not included in **argc** or **argv**, and may appear anywhere in the argument list. [The equivalent HDOS command would be

```
progl a b <syl:infile >lp: ]
```

Since the C/80 compiler is itself a C/80 program, I/O redirection may be used to redirect to a file the error messages usually output to the terminal.

9.4. Interrupting a Program:

C/80 programs may be interrupted by **ctrl-B** or **ctrl-C** [**ctrl-C** only on HDOS], which causes the program to terminate immediately and exit. The standard output is closed (except by **ctrl-C** on CP/M), but all other open output files are lost.

NOTE (applies to CP/M users only): Under CP/M, a **ctrl-C** may not always be noticed when it is typed. The C/80 library will check for a **ctrl-C** whenever a disk read or write is performed, but if a program does not access disk it must take special steps in order to be interruptable.

The HDOS operating system provides a way for the user to trap and handle interrupts caused by **ctrl-A**, **ctrl-B** and **ctrl-C**. C/80 provides a similar, but weaker, capability under CP/M. When the C/80 console input routine detects a **ctrl-B**, a call is executed to the location contained in **CtlB**. That location initially contains **exit()**, and thus **ctrl-B** usually aborts a running program, just as **ctrl-C** does.

Instead of exiting, your program may choose to handle **ctrl-B** interrupts itself. To do so, declare **CtlB** to be an **extern int**, and set it to a subroutine to be executed when **ctrl-B** is typed. If that subroutine returns when done, program execution will continue.

Since CP/M does not usually detect a typed character, even **ctrl-C**, until the program attempts to read from the keyboard, C/80 programs will often not respond immediately to a **ctrl-B** or **ctrl-C**. The C/80 library does check for these characters whenever a disk read or write is performed. You can make your program check more often by calling **CtlCk()** whenever a check is desired.

One side effect of **CtlCk()** is to read any typed character that may be waiting at the console. If the character is not a **ctrl-B** or **ctrl-C**, it will be echoed and will be placed in the input line buffer for the next call

to `getchar()`. So calling `CtlCk()` provides a limited typeahead capability. However, editing characters like `ctrl-U` and `DEL` do not operate on characters read in this manner. Also note that once the buffer has been filled (about 130 characters), any further typed characters will be echoed but discarded without warning, until the buffer has been completely emptied.

Sometimes this typeahead can be annoying. It is suppressed if `Cmode` is set to 0. It can also be suppressed by patching out `CtlCk()`. This can be done in a crude but effective fashion by declaring it a char and storing 0311 (RET instruction) into it.

9.5. Basic I/O Library Routines:

The basic I/O routines are included in the `CLIBRARY.ASM` and `CLIBRARY.REL` files, one of which must be included in any C/80 program. Your program can call these routines without doing anything special to load them.

getchar() - returns a character from the standard input (usually the terminal). -1 is returned for end of file (`ctrl-D` under HDOS; `ctrl-Z` under CP/M).

putchar(c) - writes the character `c` on the standard output (usually the terminal), and returns `c`.

fopen(fname,mode) - opens the named file and returns the channel number of the file. `Fname` is a string constant, or pointer to or array of characters containing the file name. The name may be any legal file or device, like `SYL:FOO.TXT`, or `TT:` (under HDOS; CP/M equivalents are `B:FOO.TXT` and `CON:`). `Mode` may be "r", "w" or "u" for read, write or update mode. (Update mode is treated the same as write mode but the file is not deleted before opening. [On HDOS, it is opened in update mode.] `Mode` may also be "rb", "wb" or "ub" if the file is to be treated as a binary file (see `getc`). If the file can not be opened, `fopen` returns 0. A file or device which is written on must be closed explicitly, or some or all of the file may be lost (except for file `fout`).

At most 6 files may be open at any one time. I/O buffers are allocated for three files. If more than three files are opened, `fopen` will call `sbrk` to allocate a buffer of 256 bytes for each additional file. If there is not enough memory available, the open may fail.

[HDOS: The `fopen` channel number is not necessarily the same as the system's channel.]

getc(chan) - returns the next character from the file or device open for reading on channel `chan`. Returns -1 to signify end of file. Channel 0 is always the terminal, even if I/O redirection has been used. If the file was opened in binary mode, `getc` will read every byte in the file. Files

opened in normal mode are treated as ASCII files. Under CP/M, ctrl-Z is interpreted as end of file, and newlines, which are a CR-LF pair in the file, are read in as the single character '\n'. [Under HDOS, the only special treatment is that 0 bytes are ignored.]

putc(c,chan) - writes the character **c** on the file or device opened for writing or update on channel **chan**. Channel 0 is always the terminal, even if I/O redirection has been used. In ASCII files, the conversions listed under **getc** are performed in reverse.

fclose(chan) - closes the file or device opened on channel **chan**. If a file opened for writing is not closed before the program terminates, the last block of the file may be lost. Once a channel has been closed, another file may be opened without exceeding the open file limit (see **fopen**). On CP/M, an end of file (ctrl-Z) is written onto the file at the current position unless the file was opened in binary mode.

read(chan,addr,n) - reads up to **n** bytes from channel **chan** into memory starting at address (pointer) **addr**. **N** must be a multiple of 128 [HDOS: 256]. **Getc** reads one character at a time; **read** provides an alternative method for reading many characters at once. **Read** returns the number of bytes read, which may be less than **n** if the end of file was encountered. **Read** returns 0 if an error occurred. **Read** does not perform any character conversion regardless of whether the file was opened in binary mode. **Read** and **getc** should not both be used on the same channel.

write(chan,addr,n) - writes **n** bytes from address **addr** to the file open on channel **chan**. **N** must be a multiple of 128 [HDOS: 256]. **Write** provides an alternative to **putc** for outputting many characters at a time. **Write** returns 0 if an error occurred, and the number of bytes written otherwise. **Write** and **putc** should not both be used on the same channel.

exit() - terminates the program and returns to command level. Does not close any open files except the standard output. Returning from **main** has the same effect as **exit()**.

sbrk(n) - allocates a block of **n** bytes of memory, returning the address of the first byte, or -1 if that much memory is not available. The allocated area grows upward from the end of the user program, and the stack grows downward from high memory. Subsequent calls to **sbrk** will always provide adjacent blocks of memory. Although **sbrk** will always leave about 500 bytes for stack expansion, it is still possible for the stack to grow into allocated memory (or static storage or program memory, for that matter), with undefined results. Note that **fopen** may call **sbrk** to allocate I/O buffers. See also **alloc** and **free** (Section 9.7).

9.6. Formatted Input and Output.

Formatted output is provided by `printf`, `fprintf` and `sprintf`, which are on file `PRINTF.C` and may be included in a compilation by placing the command `#include "printf.c"` before the first reference to `printf`. If `printf` is included in any other way (e.g., by `XTEXT PRINTF.ASM` or by linking `printf.rel` using `LINK` or `Link-80`), the header file `PRINTF.H` must be included instead.

Formatted output is provided by `scanf`, `fscanf` and `sscanf`, which are on file `SCANF.C`. The corresponding header file is `SCANF.H`. These routines are incorporated analogously to `printf`.

These routines provide functions similar to the routines described in The C Programming Language.

`printf(stg,v1,...,vn)` - prints the values `v1` through `vn` (`n` \geq 0) on the standard output (usually the terminal), using `stg` as the format specification. The characters in `stg` are printed on the standard output, except for conversions. Each conversion takes the next value from the argument list and prints it as the conversion specifies. A conversion consists of:

The character `%` (required).

An optional minus sign, specifying left justification of the value in the field.

An optional decimal number specifying a minimum field width. If the value is too wide for the field, the field will be expanded. If it is too narrow, the field is padded with blanks. If the first character of the field width is 0, however, the field is padded with zeros.

An optional precision, consisting of a period and a decimal number. This has meaning only for a string, and specifies the maximum number of characters to be printed.

A conversion letter (required), specifying how the value is to be printed. The conversion letters allowed are:

- `%d` (decimal number output, signed)
- `%o` (octal, unsigned)
- `%c` (single character)
- `%s` (string)
- `%x` (hexadecimal)

`fprintf(chan,stg)` - like `printf`, but output goes to the file opened on channel `chan` rather than to the standard output.

`sprintf(addr,stg)` - like `printf`, but output goes into the character array beginning at memory location `addr`.

For example, the program fragment

```
printf("i = %d, s = %s\n",27,"Hi there!");
```

would print out the line

```
i = 27, s = Hi there!
```

The file TPRINTF.C contains a more compact version of `printf` which lacks left justification and precision.

`scanf(stg,pl,...,pn)` - inputs values from the standard input into the locations pointed to by `p1` through `pn` ($n \geq 0$), using `stg` as the format specification. `Stg` contains conversions, consisting of:

Blanks, tabs and newlines, which are ignored.

Characters other than `%`, which are supposed to match identical characters in the input.

Conversions, beginning with the character `%`, an optional `*` (to suppress storing the value), an optional maximum field width, and a conversion character. The conversion characters are as shown, and should correspond to an argument which points to an object of the corresponding data type:

```
%d (decimal - int)
%o (octal, unsigned - int)
%c (single character - char)
%s (string - char array)
%x (hexadecimal - int)
%h (decimal - short int; same as int)
```

With the optional MATHPAK, the conversion characters `d`, `x` and `o` may be preceded by an `l` (letter L) to indicate that the corresponding argument is a long. Also, the conversions `e` and `f` are allowed to indicate floating point conversion.

The effect of `scanf` is to read the standard input, skipping blanks, tabs and newlines. When a nonblank character is encountered, it is taken as the start of an input field. The field is converted according to the next conversion in the format string, unless the `*` appears to specify assignment suppression. The field ends on the first blank, or, when the field width is exhausted, whichever comes first. If conversion is successful, the value is stored in the object pointed to by the next argument to `scanf`.

`Scanf` returns when the input argument list or conversion string is exhausted, or when it encounters an input that can not be converted according to the specification.

`Scanf` returns the number of items converted and stored, or -1 to indicate end of file. IMPORTANT: In order to use

the value returned by `scanf`, the `scanf` call must be enclosed in an extra set of parentheses, because of the kludge used by C/80 to permit multiple arguments.

The most common mistake in using `scanf` is to put variables in the argument list instead of pointers. This is very easy to do, and will cause random locations in memory to be overwritten. When `scanf` makes strange things happen, remember to check that each variable in the argument list (except array names and pointers) is preceded by a `&`.

`fscanf(chan,stg)` - like `scanf`, but input comes from the file opened on channel `chan` rather from the standard input.

`sscanf(addr,stg)` - like `scanf`, but input comes from the character array beginning at memory location `addr`.

9.7. More Storage and I/O Routines.

The following routines provide memory allocation, file manipulation and I/O capabilities beyond that afforded by the basic library. They are contained in the file `STDLIB.C`. This file contains `#ifneed` directives around each routine, so if you include it at the end of a program source file, C/80 will compile only those routines actually called. For Macro-80 and RMAC users, a library file `STDLIB.REL` is also provided.

`alloc(n)` - returns a pointer to `n` consecutive bytes of memory, available for use by the calling program.

`free(p)` - Returns to the free memory pool the block of memory pointed to by `p`. `P` must have been a value returned by `alloc`.

`getline(stg,len)` - Reads a line from the console into the char array pointed to by `s`. Stops after `len-1` bytes or when a newline is typed. Terminates `stg` with a 0 byte; does not store the newline. Returns the number of bytes in `stg`, excluding the terminating byte.

`rename(s,t)` [CP/M only] - Renames the file whose name is in the string `s` to have the name in the string `t`. Returns -1 for failure, a non-negative number for success. Failure occurs if file `s` does not exist or if `s` and `t` are on different devices.

`unlink(s)` [CP/M only] - Deletes the file `s`, if it exists.

9.8. Arithmetic and Number String Functions.

The following routines from `STDLIB.C` provide arithmetic capabilities and conversion between ints and strings. Most are coded in C and are provided for programming convenience, not speed.

abs(i) - Returns the absolute value of the integer i.

atoi(s) - Returns the integer value of the number contained in the string s; stops at the first non-digit.

itoa(i,s) - Converts the integer i into an ASCII string in the char array s[7]; returns s.

max(i,j) - Returns the greater of the integers i and j.

min(i,j) - Returns the lesser of the integers i and j.

9.9. String Manipulation.

The following routines from `STDLIB.C` provide string manipulation capabilities. Some are fast assembly language routines, but most are simply coded in C and provided for programming convenience.

Strings are assumed to be terminated by a 0 byte. No checking is done to determine whether copied or appended strings overflow the arrays they are placed in.

index(s,t) - Checks to see if the string s is a substring of the string t. If so, returns the starting position in t of s; if not, returns -1.

isalpha(c) - Returns 1 if c is an alphabetic character (A-Z or a-z), otherwise 0.

isdigit(c) - Returns 1 if c is an ASCII digit (0-9), otherwise 0.

islower(c) - Returns 1 if c is a lower case alphabetic character (a-z), otherwise 0.

isupper(c) - Returns 1 if c is an upper case alphabetic character (A-Z), otherwise 0.

isspace(c) - Returns 1 if c is a blank, tab or newline, otherwise 0.

strcat(s,t) - Copies string t onto the end of string s.

strcmp(s,t) - Compares string s to string t, returning -1, 0 or 1 if s is less than, equal to, or greater than t. Inequality is computed by numerical ASCII value. In particular, strings containing only upper or only lower case letters are compared in alphabetical order, but all upper case letters are less than any lower case letter.

strcpy(s,t) - Copies string t into memory starting at the pointer s.

strlen(s) - Returns the number of bytes in the string **s**, exclusive of the terminating 0 byte. Note that it takes **strlen(s)+1** bytes to hold string **s**.

tolower(c) - Returns the character **c**, but if **c** is an upper case letter (A-Z) it is converted into the corresponding lower case letter.

toupper(c) - Returns the character **c**, but if **c** is a lower case letter (a-z) it is converted into the corresponding upper case letter.

9.10. CP/M System Calls.

The following routines from **STDLIB.C** provide the ability to call the CP/M system (BDOS) to manipulate files and perform other basic operations. Consult the CP/M 2.2 Interface Guide (one of the manuals provided with CP/M) or equivalent for details on individual system calls.

[These routines will not work on the HDOS operating system.]

bdos(c,de) - Performs a CP/M BDOS system call with registers **C** and **DE** set to the values shown. Returns the contents of the **A** register as a 16 bit sign extended number.

makfcb(s,fcb) - Converts the file name **s** into a CP/M file control block (FCB) in the char array **fcb[36]**. **S** may contain an optional disk drive identifier and extension, and upper and/or lower case letters.

9.11. Random Access File I/O:

The file **SEEK.C** contains routines affording random access file I/O capability. This file can be included in your C/80 program by the statement

```
#include "seek.c"
```

(Note: CP/M 1.4 and earlier CP/M releases do not support random file I/O, and these routines will not work on those versions of CP/M.)

SEEK.C contains the following library routines:

seek(chan,offset,type) - moves to a specified position in the file which is open on channel **chan**. The next **getc** or **putc** call will read or write starting at the new location. The value **offset**, which may be positive or negative, specifies the number of bytes that the read/write pointer is to be placed from:

type = 0: the beginning of the file.
type = 1: the current read/write location.
type = 2: the end of the file.

For example, `seek(chan,0,2)` will position the read/write pointer at the end of the file. If `type = 3, 4` or `5`, the pointer is moved `offset` records (256 bytes) instead. `Seek` returns a value of `-1` if an error occurs, `0` for success.

`ftell(chan)` - returns the current read/write pointer for the file open on channel `chan`. This pointer is the number of bytes before the current position in the file. If the current position is greater than 65535, the value returned will be correct mod 256; i.e., the byte position in the current 256 byte record will be correct, but not the record number.

`ftellr(chan)` - returns the current read/write pointer for the file open on channel `chan`, divided by 256.

It is possible to both read and write to the same file without reopening it, if the file has been opened in write or update mode. However, remember that closing a file opened for write or update will result in an end of file character being written at the current position, unless it was opened in binary mode. So if you have written into the middle of an ASCII file, you should seek to the end before closing it. You may prefer to use binary mode and supply your own end of file characters. But if you do so, remember that in binary, CP/M end of line is indicated by the pair `"\r\n"`.

9.12. Program Chaining; Wildcards

The file `EXEC.C` contains a routine which allows execution of another program from within a C/80 program.

`exec(prog,args)` - chain to another program. `Prog` is a string containing the name of a program. `Args` is a string containing any arguments, separated by blanks, just as on the command line. `Exec` will execute the named program with the arguments given, just as if it had been invoked from the command line. Unless an error occurs, control never returns from `exec`. All open files must be explicitly closed before calling `exec`, or strange things may happen.

The file `COMMAND.C` contains a routine which allows UNIX-style command line expansion of file names using wildcards.

`command(&argc,&argv)` - expand file name wildcards in the command line. `Argc` and `argv` are the arguments to `main`. `Command` will treat any element of `argv` which contains the characters `'?'` or `'*'` as an ambiguous file name, and will replace it by the zero or more strings containing the names of files which match the specification. It calls the routines `bdos` and `makfcb`, which are also on `STDLIB.C`.

10. USING C/80 WITH MACRO-80 OR RMAC

C/80 can optionally generate assembly code for input to the Microsoft Macro-80 or Digital Research RMAC relocatable assemblers. This allows you to develop a C/80 program in several modules, generating .REL files which can be linked using the LINK or Link-80 linking loader.

You can create .REL modules for C/80 library routines such as PRINTF and SEEK, to speed up their inclusion in your programs. A library manager such as LIB-80 can be used to create a library containing both your commonly used subroutines, and the individual functions in the C/80 library. Such a library can be selectively loaded using the S switch of the linker.

Developing a large program in this way is more efficient both in time and disk space. Although AS is sufficient for learning C and for much serious work, use of a relocatable assembler is recommended for large projects.

10.1. Assembling and Loading.

To generate a .MAC file, invoke C/80 using the -m switch:

```
c -m [other args ...]
```

To generate an .ASM file compatible with RMAC, use -m2 instead. You can also configure C/80 so that one of these modes is the default. See Section 6.2.

Macro files generated by C/80 must be assembled by M80 or RMAC to create a .REL file. REL files are then linked by L80 or LINK to create a .COM file [HDOS: .ABS file]. IMPORTANT: When linking C/80 .REL files, the file CLIBRARY.REL must be linked in and it must be the last file linked.

Using a relocatable assembler imposes a few additional restrictions on global identifiers. In global variable and function names, only the first six characters are significant (as opposed to 7 when AS is used), and upper and lower case are considered identical. Global arrays and variables must be defined in only one source module, and must be declared **extern** by any other module that references them, being careful to distinguish between arrays and pointers. Functions are implicitly **extern** and need not be declared.

Macro-80 will not allow the use of a register name as an identifier. Thus, you may get assembly errors when you use the following as function names or globals, in either upper or lower case:

A, B, C, D, E, H, L, PSW, HL, DE, BC

RMAC also uses opcodes as identifiers, so it imposes many more restricted names. Generally, it is all right not to worry about this restriction when writing your programs, as long as you

recognize what is happening when the assembly errors occur.

Since RMAC does not accept the '_' character in identifiers, C/80 substitutes '?' for it when RMAC compatible output is selected.

C/80 produces CSEG and DSEG directives to load code and data into separate areas of memory. This has the side effect of requiring more memory for the linking loader. If the linker should run out of memory, try compiling with the -a switch, which suppresses generation of CSEG and DSEG.

10.2. ROMable Code.

C/80 may be used with a relocatable assembler in order to generate read-only code for insertion into read-only memory (ROM). The compiler generates CSEG and DSEG directives to load program code and data into separate areas of memory.

Portions of the C/80 runtime library are not read-only, however, and in addition depend on the operating system environment which is often not available in ROM-based applications. For this reason, the source code for this portion of the runtime library is furnished in the file CLIBIO.C on the C/80 distribution disks.

The file CLIBRARY.REL consists of the relocatable modules from CLIBIO.C, the arithmetic library CLIBMATH, and a small module .FIXMSOFT, defining three 7-character library entry point names for compatibility with Microsoft Macro-80 version 3.43, which for some reason generates 7-character globals. Lib-80 can be used to extract these modules for use in generating a ROMable library.

10.3. Making Libraries.

You will probably want to create a single large library file which you can search on every load. If you put all the C/80 library functions into it, as well as your own private library functions, you will not need to worry about including the functions called by your program, as this will happen automatically.

In creating such a library, you must be careful about the order of the modules you include. If a function in one module calls a function in another module, the caller must come before the called function, so that the linker can load them both without backing up, something it will not do.

As an example, here are the commands you would type to use LIB-80 to create a library CLIB.REL which includes the functions in STDLIB, PRINTF and CLIBRARY.

```
A>lib
*clob=stdlib,printf,clibrary
*/e
A>
```

11. MULTIPLE COMPILES USING AS.

To reduce space and time taken by compilations, it is often useful to compile a large program in several pieces. This may be done by splitting the program into several C source files, using the `extern` declaration to reference globals which are declared in another source file. For example, if the declaration `int i[5];` appears at the top (global) level in one source file, the declaration `extern int i[5];` in a second source file will allow programs in the second file to refer to the array `i` defined in the first file. Simply saying `int i[5]` in both files will cause `i` to be doubly defined at assembly or load time. Exception: functions defined in one source file may be called from another source file in the same assembly or load without any special declaration.

There are two ways to generate a single object module from multiple compilations. One way is to use Macro-80 and Link-80 from Microsoft, or RMAC and LINK from Digital Research, as described in the previous section. The other is to use the AS assembler provided with C/80 to assemble one module from several .ASM files.

An example will illustrate how to do this. Suppose there are three C source files: MAIN.C, SUB1.C, and SUB2.C. First you must insert into MAIN.C statements which will cause the assembly language files for SUB1 and SUB2 to be included in the assembly. The following statements should be put into MAIN.C somewhere at top level (i.e., not inside a subroutine):

```
#asm
      XTEXT SUB1.ASM
      XTEXT SUB2.ASM
#endasm
```

The XTEXT command will cause these lines to be included in the assembly language file MAIN.ASM.

Next, compile MAIN.C to create an output file called MAIN.ASM. This file contains compiler-generated labels, which usually look like .a, .b, and so on.

Then compile SUB1.C to produce an output file named SUB1.ASM. If nothing is done to prevent it, the compiler will use the same labels .a, .b, etc., in SUB1.ASM, and there will be a conflict when the two files are assembled together. To prevent this, SUB1.C should be compiled with a command like "C -L1000", which will start the labels 1000 down in the sequence. (The largest permissible value for the -L switch is 32767). The -L switch also suppresses the generation of instructions in SUB1.ASM to include the C I/O library, since MAIN.ASM already has those instructions.

Similarly, compile SUB2.C, say by "C -L2000", to produce SUB2.ASM. If disk space is a problem, these compilations may all be performed on different disks, and the .ASM files copied to another disk for assembly. The library file CLIBRARY.ASM must reside on A: [SY0: on HDOS] during assembly.

To assemble the program, give the command `AS MAIN`. All the

files will be assembled to produce MAIN.COM [MAIN.ABS on HDOS] which can then be run.

12. RUNTIME TRACE AND EXECUTION PROFILE

Most programs which take very long to run spend most of their execution time executing a relatively small amount of code. C/80 contains a runtime execution profile feature to help identify where a program is spending its time, so the critical routines can be made more efficient.

To use this feature, compile the program with the `-p` switch, as in

```
c -p progname
```

If you are using AS, you will need to create the file CPROF.ASM and make sure it is on A: [SY0: on HDOS] when the program is assembled. To generate CPROF.ASM from CPROF.C, use the command

```
c -L32600 cprof
```

If you are using Macro-80 or RMAC, compile and assemble CPROF.C to produce CPROF.REL, and link it in when you load.

Now run the program. When the program finishes running and exits normally, a listing will be produced on the standard output device showing, for each subroutine, the number of times it was called and the total time (in units of two ticks of the computer clock) spent inside the subroutine. You can save the output listing by using I/O redirection to send the standard output to a file instead of the terminal.

Similarly, a runtime trace of the program execution can be produced by copying file CTRACE.ASM from the distribution disk onto A: [SY0: on HDOS] as file CPROF.ASM, and compiling as above.

Note (applies to CP/M users only): C/80 can only provide execution times if your system has a 16-bit clock at some address in memory. The Heath/Zenith systems have a clock at memory location 0B hex. To use the profile feature on systems with no clock, or with a clock at another location, you must regenerate file CPROF.ASM as follows:

Edit file CPROF.C. Locate the line which begins with `#define TICCNT`. If your system has a clock, replace the expression following `TICCNT` with the address of the clock word in memory. If there is no clock, remove the entire line. Then recompile CPROF by the command

```
c -L32600 cprof
```

13. THE AS ASSEMBLER.

The C/80 distribution includes AS.COM, an absolute 8080 assembler. AS is essentially the same as the ASM assembler under Heath HDOS. This section gives a brief description of AS, to help you write assembly language code to be included in C/80 programs.

To assemble a program, type a command of the form

as comfile,listfile=infile

where comfile is the name of the absolute file to be produced, listfile is the file on which to write an assembly listing, and infile is the assembler source file. If no extensions are specified, they default to .COM, .LST and .ASM respectively. Listfile can be a C/80 device, such as LST:; if that device does not respond to tabs, it is more useful to list to CON: and use ctrl-P to obtain a hard copy. The command

as filename

is short for

as filename=filename

AS takes Intel 8080 mnemonics, upper case only. Identifiers are up to 7 characters from the set A-Z, a-z, 0-9, ., _, ?, @ and \$. Upper and lower case letters are distinguished in identifiers. For compatibility with other assemblers, such as Macro-80, colons following a label are ignored, as are '#' characters following a label in an expression.

Constants can be one or two characters enclosed in single quotes, or a string of digits, possibly with a suffix O or Q for octal, H for hex, or B for binary. Default is decimal.

The symbols * and \$ represent the address of the current instruction.

Arithmetic expressions in an address field are evaluated strictly left to right. The operators are +, -, *, /, & (bitwise and), and < (left shift). Parentheses are not allowed.

The following pseudo-ops are identical to the ones in the CP/M ASM assembler: ORG, EQU, DB, DW, DS.

The pseudo-op "XTEXT filename" is the AS equivalent of the C/80 #include preprocessor directive. It includes the named file at that point in the assembly. If no disk is specified, it assumes the disk on which the current source file resides.

The pseudo-ops "LON ccc" and "LOF ccc" control listing. "ccc" is a string of characters: L turns listing on (LON) or off (LOF), C controls listing of lines from XTEXT files (default off), and G lists all bytes generated by an instruction (default: list just the first five).

14. UPPER CASE SOURCE FILES

Since the C/80 language depends heavily on the full ASCII character set and on lower case keywords, a facility is provided to allow C/80 source files to be prepared on upper case only terminals. Upper case source files should contain the preprocessor command

#UPPER

as the first line in the file. This causes the compiler to interpret the remainder of the file in upper case mode. In this mode, the compiler translates each upper case letter to the corresponding lower case letter. Upper case letters, and the special characters not available in the upper case ASCII subset, are typed as follows:

The characters:are interpreted as:

```

\ (
\ )
\ !
\ -
\ _
A ... Z
^A ... ^Z

```

```

{
}
~
^
a ... z
A ... Z

```

In upper case mode, when the character ^ appears before a non-alphabetic character, it is ignored. Note that ^ may be displayed on some terminals and printers as an up arrow, and on others as a caret or "hat".

The #UPPER command affects only the file in which it appears, and does not have any effect on an #include file. If an #include file is prepared in upper case mode, it must contain its own #UPPER command. (Thus, PRINTF.C, which is not in upper case mode, may be included in a file which is in upper case mode.)

15. TRICKS AND INTERNALS

15.1. Global Arrays.

Global arrays occupy space in the .COM [or .ABS] file. To reduce the size of these files, it is best to allocate large arrays at run time. This can be done by making them local to a function, or by using `alloc()` or `sbrk()`.

15.2. Assembly Language Linkage.

The C function call linkage is to push the arguments, as 16-bit quantities, on the stack, and call the subroutine. It is the responsibility of the calling program to pop the arguments off the stack on return. Note that standard C pushes its arguments in reverse order; this means a C/80 routine may not usefully be called with a different number of arguments than it expects.

The value of a function is returned in the HL register.

If 32-bit quantities are used with the optional MATHPAK, those values are pushed on the stack high word first, and values are returned in the BCDE registers.

Assembly language routines invoked by C function calls have access to all registers and need not restore their values.

The compiler generates code using a primary register (HL), a secondary register (DE), and the stack. Most operations are performed by calling library subroutines, with the first operand being HL and the second operand either DE or the top of the stack (which is usually popped as a side effect of the function).

If you want to fetch the value of a variable using in-line assembly language code, put a statement consisting only of the variable just before the `#asm` directive. C/80 will leave the variable value in HL.

15.3. Declarations for Efficient Code.

In order to produce code which is both space and time efficient, variables should be declared **static** whenever possible, and **int** in preference to **char**. Declaring a function argument to be **register** will save space if the argument is used five or more times in the body of the function, and will save time if the argument is accessed often during function execution.

15.4. Defining Globals in Header Files.

When compiling using several source modules, it is handy to

define global variables using an included header file. But these variables must be defined only in one module, and declared **extern** in all the others. This can be accomplished through the following programming trick:

In the header file, declare all the variables as follows:

```
#ifndef EXTERN
#define EXTERN extern
#endif
EXTERN int i, j, ...
```

Then, in the one file in which the variables are to be defined, place

```
#define EXTERN
#include "header-file"
```

15.5. I/O Buffers.

At runtime, the library allocates three I/O buffers directly below the operating system area, and then builds the stack downward starting below the buffers.

Note (applies to HDOS users only): Under HDOS, device drivers may load below the overlay area when the device is opened and closed. If a device driver spills over the 768 byte buffer area and into the stack, problems may occur. Since HDOS provides no satisfactory method for determining potential driver memory requirements, the runtime library can not anticipate this problem automatically.

The problem will rarely arise, since most systems have at most a single LP: or AT: device (in addition to TT: and SY:, which are always loaded), and no single driver is large enough at present to cause trouble. If multiple device drivers are to be loaded it may be necessary to leave additional room for the drivers. To do this, increase the number 10 in the LXI B,10 instruction in CLIBRARY.ASM (around line 40), adding the number of bytes necessary (1000 per extra device driver should be generous).

15.6. Crash on Exit.

A common condition on CP/M is for a program to run to completion, and then crash during exit. The system can go completely dead, or a BDOS error message can complain about a nonexistent condition or drive.

This is often caused by storing into location 0, through a pointer whose value is initially 0 and is never set. C/80 programs exit via a jump to location 0. (Exiting via system call 0 would have been safer, but this is known to conflict with DESPOOL, although both it and CP/M are from the same manufacturer!)

16. COMPILER ERROR MESSAGES

When a C/80 program contains a detectable error, the compiler will produce an error message, giving the source file name and line number, and a description of what the compiler thinks the error is. The source line is listed, with an arrow pointing to the location of the error in the line. (The line is shown the way it looks after all `#defines` have been expanded.)

These messages aren't always as helpful as one would like, however. For instance, the compiler may be looking for a different statement type than the one you thought you wrote. So the message may not describe the error. Sometimes, the compiler detects the error far away from the place where it actually occurred. For example, leaving off a `'}'` deep inside a function will cause the next to last `'}'` in the function definition to terminate it, at which point the compiler will probably spew out dozens of error messages as it tries to parse executable statements as declarations. To find such an error, you may have to inspect many lines of code.

In addition, once an error has been detected, the compiler is not always able to recover and continue parsing the remainder of the program. So a single error will sometimes result in a large number of error messages. Often, this is the result of C/80 trying to parse statements as declarations or declarations as statements. When you can't understand what some of these messages are complaining about, you should fix the first error detected, and any others you can easily locate, and then recompile to find any remaining ones.

Limitations on error detection and analysis are almost unavoidable in compilers, especially when trying to cram a powerful language into the limited space of a microcomputer. This section can help by describing C/80's error messages and what might cause them to occur.

bad label: Labels must follow the rules for identifier names.

bad type: C/80 expected to see a type: short, long, char, int, unsigned, struct or union.

can't initialize auto: You can only initialize globals or local statics.

can't initialize union: Like it says.

can't subscript: You can only subscript a pointer or an array. (Did you declare it?)

can't find file: Did you specify the file and extension correctly? Did you leave off the device? Also, when available memory is almost exhausted, the compiler may be unable to allocate the buffer space to open an `#include` file.

can't compute size: Maybe the arrays in the expression have been declared incorrectly.

construct not permitted: C/80 does not allow nested type declarations. For example, **sizeof** an abstract type can not be used as a dimension in a declaration.

dimension missing: The only time an array dimension can be omitted (or 0) is when declaring a function argument, or when the dimension is determined by the size of an initializer. Furthermore, only the last dimension of an array can ever be omitted.

extra ; (ignored): This looks like a function declaration, except function declarations are `f(...) { ...}` and you put a `;` after the `)`. So the compiler took it out.

ifdefs nested too deep: Maximum nesting of `#ifdef`, `#ifndef`, `#if` and `#ifneed` is 5.

illegal constant value: Cases must be int or char constant values.

illegal expression - need lvalue: Some operations (`&` and assignment, for example) require an lvalue, which is an object with a memory address. This isn't one. (Things that aren't lvalues include constants and expressions.)

illegal function call: The identifier or expression is not of type function.

illegal initializer: The initializer expression can not be computed at compile time. It must be an int or char constant expression.

illegal struct reference: Either a `.` preceded by something that isn't a structure, or a `.` or `->` followed by something that isn't a structure element. Remember that `.` is used following things that are structures, and `->` following things that point to structures.

illegal symbol name: The compiler wanted an identifier here.

improper argument: This argument is too large; probably a structure or union. Try using a pointer to it instead.

internal compiler error: The compiler encountered an error in code generation. If there are previous errors, fix them first. If not, this may indicate a compiler bug; please report it to The Software Toolworks. You may be able to proceed by simplifying or rearranging the expression.

invalid expression: The compiler is looking for an expression, but this does not look like one.

line too long: Input line longer than about 100 characters.

macro table full: Recompile using the `-d` switch to increase the size of the `#define` table; see Section 6.1.

misplaced case: case not inside a switch statement.

- ? missing:** (where ? is some punctuation character): C/80 expected to find that character and didn't. It inserted it and continued, so if you really did leave that character out at that spot, the compilation proceeded correctly.
- must be a constant:** The compiler looked for a constant or constant expression, but didn't find it. Remember that ?: is not legal in a C/80 constant expression.
- nested too deep:** #include files can be nested no more than three deep (a main file and two included ones).
- nesting:** Structs and unions can not contain an instance of themselves.
- no active whiles:** A break or continue statement was not inside any for, while or switch statement.
- not a label:** This construct needs an identifier which is a label. You have used something which is either not an identifier or something besides a label (like a char or int).
- not a declared variable:** The compiler wants an lvalue here. (See "illegal expression - need lvalue" above.) Usually this error means you have forgotten to declare an identifier; all identifiers in C must be declared before use.
- not a function:** This identifier or expression is followed by a '(', so it looks like a function call. But it's not of type function.
- not an address:** In an initializer, you applied & to something that does not have an address.
- not a pointer:** You applied * to something that is not a pointer. (Maybe you applied it implicitly, like in a subscript operation.)
- not enough memory for long line (fatal error):** Through use of #defines, you have created an input line which can not fit into memory.
- operands and/or operator incompatible:** This usually means you have performed an illegal arithmetic operation on a pointer. The only legal ones are pointer plus integer, and pointer minus either an integer, or another pointer that points to something of the same size. If the expression looks legal, try rearranging the order of the operands. If you know what you are doing but C won't let you (like computing pointer & mask), use the (int) cast to fake the compiler out.
- output file error:** Usually means the disk is full. You won't see this for long, because the compiler will start dumping the assembly language output to the terminal.
- previously defined:** This identifier has been defined before. Externs, in particular function names, may be redeclared

as long as the type is not changed. However, if a previously undeclared function is called, it is implicitly declared **extern int**. If it is redeclared or defined as **static**, or of a different type, you will get this error. Avoid it by declaring the function's correct type before it is first called.

string space exhausted: Recompile using the **-c** flag to increase the size of the string table, or use the **-k** flag to avoid storing strings from the entire compilation. (see Section 6.1).

struct table overflow: Recompile using the **-r** flag to increase the size of the structure table (see Section 6.1).

struct too large: Only 50 elements are allowed in a struct.

symbol table overflow: Recompile using the **-s** flag to increase the number of symbol table slots allocated (see Section 6.1).

syntax error: The compiler was trying to find a declaration (at the top level) but couldn't.

too complex: In an initializer, this expression was too complicated, or was not a constant.

too complicated: A type declaration had more than 7 levels of indirection (*****, **[]** or **()**).

too large for register: Register declarations may only be applied to chars, ints and pointers.

too long: The initializer string, even minus the terminating 0 byte, exceeds the size of the char array.

too many active whiles: Well, congratulations. There's only one table in this compiler that isn't expandable, and you have overflowed it by nesting 20 **whiles**, **fors** and/or **switches**. Simplify your program.

too many cases: Recompile using the **-w** switch to expand the switch case table (Section 6.1).

too many structs: At most 239 different struct types may be declared.

type mismatch: You tried to initialize an identifier with a value of the wrong type.

undefined struct name: This struct type has not been declared.

usage error: Indicates a violation of some C usage rule, like passing a struct as a function argument.

warning: =? op assumed: You are using the old style assignment operator (like **=&**), and failed to leave a space between the operator and the following operand. The compiler assumes you mean **=&**, but wants you to make sure.

INDEX

abs.....31
 acknowledgements.....2
 additional reading.....5
alloc.....30
 anachronisms.....22
argc, argv.....24
 arithmetic functions.....30
 arrays.....15
 AS assembler.....7, 38
#asm.....21
 assembly errors.....34
 assembly language.....21
 assembly linkages.....40
atoi.....31
auto.....17, 21

bdos.....32
 binary file.....26
 blanks.....20
 blocks.....21
 books on C.....5
break.....19
 buffers, I/O.....21, 41
 C library files.....22
 C manual.....5
case.....19
 case conversion.....31
 CCONFIG program.....8, 12
 chaining.....33
 changing compiler defaults.....12
 channels.....26
 character constants.....17
 character I/O.....26
clibio.c.....8
clibrary.rel.....8
command.....33
 command line expansion.....33
 comparing strings.....31
 compiler defaults, changing.....12
 compiler, how to run.....10
 compiler switches.....10
 compound statement.....21
 concatenating strings.....31
 configuration program.....8
 console I/O.....23

constant expression.....17, 22
 constants.....17, 20
continue.....19
 conversion of cases.....31
 copying strings.....31
cprof.c.....8
 crash during exit.....41
CtrlCk.....25
 ctrl-B.....25
 ctrl-C.....25

 data types.....13, 15
 declarations.....21
 decrement operators.....18
#define.....22
 deleting file.....30
 developing large programs.....34
 device names.....23
 distribution disk.....7
do.....19

 efficiency.....40
 end of file.....23
 end of line mappings.....23, 26
#endasm.....21
EOF.....22
 error messages.....42
 errors, assembly.....34
exec.....33
exit.....25, 27
 exit, crash during.....41
 expression, constant.....22
 expressions.....13, 18
extern.....15, 17, 21, 45

 false.....18
fclose.....27
FILE.....22
 file, binary.....26
 file control block.....32
 file, deleting.....30
 file mode.....26
 file names.....24
 file, renaming.....30
 files, limit on open.....26

fin	24	Kernighan and Ritchie	5
fopen	26	language restrictions	20
for	19	large program development	34
formatted I/O	28	length of strings	32
fout	24	lexical conventions	20
fprintf	22, 28	libraries, making	35
free	30	library	7
free storage	30	library files in C	22
fscanf	30	library, runtime	23
ftell	33	limit on open files	26
ftellr	33	#line	22
function call linkage	40	LINK	34
getc	26	Link-80	34
getchar	24, 26	local	15
getline	30	logical device names	23
global	15	machine requirements	3
HDOS	4	Macro-80	11, 34
hexadecimal constants	17	main	24
hints	40	makfcb	32
identifiers	20	making libraries	35
identifiers, restrictions on ..	34	manual, C	5
#ifneed	22	max	31
implementation dependencies ..	20	memory allocation	27, 30
increment operators	18	min	31
index	31	multiple source files ..	34, 36, 40
initialization	21	names	15, 20
initialization to zeros	12	names, restrictions on	34
in-line assembly language	21	new features in 3.0	6
input, formatted	28	NULL	22
internals	40	number to string conversion ..	30
interrupting program	25	octal constants	17
I/O	22	operators	14, 18
I/O buffers	26, 41	output, formatted	28
I/O, console	23	pointers	16
I/O, random file	32	preprocessor commands	15
I/O redirection	24	printf	22, 28
isalpha	31	printf.c	7
isdigit	31	printf.h	7
islower	31	profile, runtime	11, 37
isspace	31	program chaining	33
isupper	31	putc	27
itoa	31	putchar	24, 26

random file I/O.....	32	string comparison.....	31
read	27	string concatenation.....	31
reading console.....	30	string constants.....	17
reading terminal.....	30	string copying.....	31
register	17, 21	string length.....	32
relocatable code.....	34	string lookup.....	31
rename	30	string manipulation.....	31
renaming file.....	30	string to number conversion..	30
restrictions on language.....	20	strings.....	15
restrictions on names.....	34	strlen	32
returned function value.....	40	structures.....	16
RMAC.....	11, 34	switch	19
ROMable code.....	35	switches, compiler.....	10
runtime library.....	23	system calls.....	32
runtime profile.....	11, 37		
		terminal, reading.....	30
sbrk	27	tolower	32
scanf	22, 29	toupper	32
scanf.c	7	tprintf.c	7
scanf.h	7	trace.....	37
scope of variables.....	15	tricks.....	40
seek	32	true.....	18
seek.c	8	truth values.....	18
sprintf	22, 28	typeahead.....	25
sscanf	30	typedef	21
standard input.....	24		
standard output.....	24	unimplemented features.....	20
statement, compound.....	21	unlink	30
statement types.....	14	#UPPER	39
statements.....	19	upper case source.....	39
static	17, 45		
stderr	22	variables.....	15
stdin	22	variables, scope of.....	15
STDLIB.C.....	30		
stdlib.rel.....	8	while	19
stdout	22	wild cards.....	33
storage classes.....	15, 17, 21	write	27
strcat	31		
strcmp	31	zeros, initialization to.....	12
strcpy	31		